

Data Engineering

TP2 :Introduction to Spark and RDDs

Diplôme National d'Ingénieur en Informatique

Spécialité :

Génie Logiciel

Réalisée par:

Oussama Ben Slama

Année Universitaire 2024/2025

Chapter 1

Spark RDD


1.1 Examples

Reading the ratings.dat file into an RDD and displaying the first 5 lines.

```
[1]: from pyspark import SparkContext
      from pyspark.sql import SparkSession
      sc=SparkContext()

      ratings=sc.textFile('ratings.dat')
      ratings.take(5)|
```

DAG Visualization :



Parsing the ratings data using parseRatings function.
 Counting how many ratings have a score of 1.
 Counting the number of unique movies rated.

```
def parseRatings(row):
    splitted = list(row.split(':'))
    return (int(splitted[0]),int(splitted[1]),int(splitted[2]),splitted[3])

ratings = ratings.map(parseRatings)


ratings.take(5)

rating_1_count = ratings.filter(lambda x: x[2] == 1).count()
rating_1_count

unique_movies = ratings.map(lambda x: x[1]).distinct().count()
unique_movies
```

DAG Visualization :

- ▶ Event Timeline
- ▼ DAG Visualization



Finding the user who rated the most movies.
 Filtering and retrieving movies rated by that user.

```
[ ]: user_id = ratings.map(lambda x : (x[0],1)).reduceByKey(lambda x,y: x+y).max(lambda x : x[1])
user_id

[ ]: movies Rated by user = ratings.filter(lambda x: x[0] == user_id[0])

[ ]: movies Rated by user.take(5)
```

Extracting and flattening all genres from the movies dataset by splitting the genre string using |.

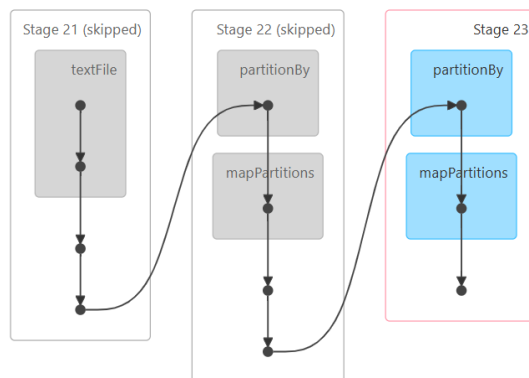
```
[18]: genders = movies.map(lambda x : x[2]).flatMap(lambda x : x.split('|'))

[19]: genders0 = movies.map(lambda x : x[2]).map(lambda x : x.split('|'))

[20]: genders.take(10)

[20]: ['Animation',
      "Children's",
      'Comedy',
      'Adventure',
      "Children's",
      'Fantasy',
      'Comedy',
      'Romance',
      'Comedy',
      'Drama']
```

DAG Visualization :




Counting the number of movies for each genre.

```
[22]: nb_genders = movies.flatMap(lambda x: x[2].split('|')).map(lambda g: (g, 1)).reduceByKey(lambda x, y: x + y)
      nb_genders.take(5)

[22]: [("Children's", 251),
      ('Fantasy', 68),
      ('Romance', 471),
      ('Drama', 1603),
      ('Action', 503)]
```

DAG Visualization :

- Event Timeline
- ▼ DAG Visualization



Get the most rated genres

```
[23]: valid_users = users.filter(lambda x : x[1] == 'M' and x[2] > 45 ).map(lambda x : (x[0],1))
      valid_ratings = ratings.filter(lambda x : x[2] >= 4).map(lambda x : (x[0] , x[1]))

      valid_rating_by_users = valid_users.join(valid_ratings).map(lambda x: (x[1][1],1))

[24]: movies_genres = valid_rating_by_users.join(movies.map(lambda x: (x[0], x[2]))).flatMap(lambda x: x[1][1].split('|'))
      movies_genres.distinct().sortBy(lambda x:x[0],1).take(10)

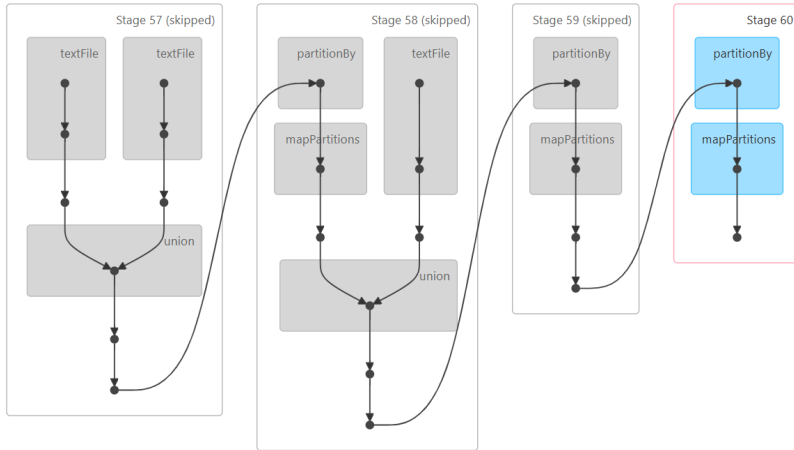
[24]: ['Action',
      'Animation',
      'Adventure',
      "Children's",
      'Comedy',
      'Crime',
      'Drama',
      'Documentary',
      'Film-Noir',
      'Fantasy']

[25]: genre_counts = movies_genres.map(lambda genre: (genre, 1)).reduceByKey(lambda x, y: x + y)
      genre_counts.take(5)

[25]: [('Action', 11988),
      ('Sci-Fi', 7137),
      ('Western', 2024),
      ('Animation', 1361),
      ('Thriller', 9586)]
```

DAG Visualization :

► Event Timeline
▼ DAG Visualization



Counting the number of movies per genre for each year

```
26]: import re
def get_year(title):
    match = re.search(r"\\(\\d{4})\\)", title)
    return int(match[1]) if match else None

27]: movies_selected = movies.map(lambda x: (get_year(x[1]), x[2])) \
    .flatMap(lambda x: [(x[0], genre), 1] for genre in x[1].split('|'))

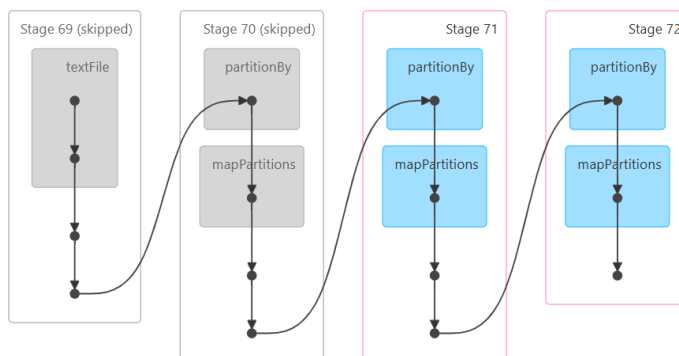
28]: movies_selected = movies_selected.reduceByKey(lambda x,y : x+y).map(lambda x: (x[0][0], (x[0][1], x[1])))

29]: movies_selected.take(10)
```

```
29]: [(1995, ('Animation', 8)),
      (1995, ('Comedy', 89)),
      (1995, ('Adventure', 25)),
      (1995, ('Crime', 18)),
      (1995, ('War', 12)),
      (1994, ('Action', 32)),
      (1994, ('Drama', 121)),
      (1994, ('Thriller', 31)),
      (1994, ('Romance', 37)),
      (1995, ('Mystery', 8))]
```

DAG Visualization :

► Event Timeline
▼ DAG Visualization



Finding the most frequent genre for each year

```
[30]: movies_selected = movies_selected.reduceByKey(lambda a, b: a if a[1] > b[1] else b)
```

```
[31]: movies_selected.sortBy(lambda x : x[0] , 0 ).take(10)
```

```
[31]: [(2000, ('Comedy', 69)),  
(1999, ('Drama', 130)),  
(1998, ('Drama', 166)),  
(1997, ('Drama', 139)),  
(1996, ('Drama', 150)),  
(1995, ('Drama', 158)),  
(1994, ('Drama', 121)),  
(1993, ('Drama', 81)),  
(1992, ('Drama', 38)),  
(1991, ('Drama', 26))]
```

DAG Visualization :

► Event Timeline
▼ DAG Visualization

